



1. Create an S3 bucket and upload some objects to S3.

Implementation


- ➔ Log in to the AWS Management Console.
- ➔ Open S3 → General purpose buckets.
- ➔ Click Create bucket.
- ➔ Enter a globally unique bucket name.
- ➔ Select the required AWS Region.
- ➔ Keep Block all public access enabled unless public access is specifically required.
- ➔ Keep Bucket Versioning disabled for this basic task.
- ➔ Click Create bucket.
- ➔ Open the newly created bucket.
- ➔ Click Upload.
- ➔ Click Add files or Add folder.
- ➔ Select the required files.
- ➔ Click Upload.
- ➔ Verify that the uploaded objects are visible under the Objects tab



2. Deploy a static website in the S3 bucket.

Implementation

- ➔ Create an S3 bucket.
- ➔ Upload the website files, for example: index.html error.html
- ➔ Open the bucket.
- ➔ Go to Properties.
- ➔ Find Static website hosting.
- ➔ Click Edit.
- ➔ Enable Static website hosting.
- ➔ Select Host a static website.
- ➔ Enter: Index document: index.html Error document: error.html
- ➔ Save the changes.
- ➔ Go to the bucket's Permissions.
- ➔ Configure the required public-access settings if you intend to expose the website directly through the S3 website endpoint.
- ➔ Add an appropriate bucket policy that allows s3:GetObject for the website objects.
- ➔ Return to Properties → Static website hosting.
- ➔ Copy the Bucket website endpoint.
- ➔ Open the endpoint in a browser.
- ➔ Verify that the static website is displayed.



Successfully edited static website hosting.

Requester pays
Disabled

Static website hosting

Use this bucket to host a website or redirect requests. [Learn more](#)

We recommend using AWS Amplify Hosting for static website hosting
Deploy a fast, secure, and reliable website quickly with AWS Amplify Hosting. [Learn more about Amplify Hosting](#) or [View your existing Amplify apps](#) [Create Amplify app](#)

S3 static website hosting
Enabled

Hosting type
Bucket hosting

Bucket website endpoint
When you configure your bucket as a static website, the website is available at the AWS Region-specific website endpoint of the bucket. [Learn more](#)
<http://poojith-1004-test1.s3-website-us-east-1.amazonaws.com>

Successfully edited bucket policy.

Bucket policy

The bucket policy, written in JSON, provides access to the objects stored in the bucket. Bucket policies don't apply to objects owned by other accounts. [Learn more](#)

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicReadForWebsite",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3::poojith-1004-test1/*"
    }
  ]
}
```

← → ↻ ⚠ Not secure poojith-1004-test1.s3-website-us-east-1.amazonaws.com

Hi, welcome to AWS S3.

← → ↻ ⚠ Not secure poojith-1004-test1.s3-website-us-east-1.amazonaws.com

Sorry, This site is under Maintainance.

3. Enable cross-region replication on S3 buckets.

Implementation:

- ➔ Create a source S3 bucket in Region 1.
- ➔ Create a destination S3 bucket in a different AWS Region.
- ➔ Enable Versioning on both buckets.
- ➔ Open the source bucket.
- ➔ Go to Management → Replication rules.
- ➔ Click Create replication rule.
- ➔ Enter a rule name, for example: CrossRegionReplication
- ➔ Choose whether the rule applies to:
 - 1. All objects, or 2. Objects matching a specific prefix/tags.
- ➔ Select the destination bucket in the other Region.
- ➔ Allow AWS to create/use the required IAM replication role.
- ➔ Save the replication rule.
- ➔ Upload a new object to the source bucket.
- ➔ Open the destination bucket.
- ➔ Verify that the object has been replicated.
- ➔ Test by uploading additional objects and confirming they appear in the destination bucket

This screenshot shows the 'Create replication rule' page in the AWS console. The 'Source bucket' section is filled out with 'Source bucket name' as 'poojith-1004-test1' and 'Source Region' as 'US East (N. Virginia) us-east-1'. Under 'Choose a rule scope', the option 'Apply to all objects in the bucket' is selected. The 'Destination' section has 'Choose a bucket in this account' selected. The 'Bucket name' field contains 'account-b-devops-bucket-050910', and a 'Browse S3' button is visible. The 'Destination Region' is set to 'US East (Ohio) us-east-2'.

This screenshot shows the 'Replication rules' page in the AWS console. It includes a warning banner about replication costs. The 'Replication configuration settings' section shows the 'Source bucket' as 'poojith-1004-test1' and 'Source Region' as 'US East (N. Virginia) us-east-1'. An 'IAM role' link points to 's3rrr_role_for_poojith-1004-test1'. At the bottom, the 'Replication rules (1)' section shows a table with one rule. The table has columns for 'Name', 'Source bucket', 'Source region', 'Destination bucket', 'Destination region', 'Status', and 'Actions'. The rule listed is 'CrossRegionReplication' with source 'poojith-1004-test1' in 'us-east-1' and destination 'account-b-devops-bucket-050910' in 'us-east-2'. The status is 'KMS-encrypted'.

Name	Source bucket	Source region	Destination bucket	Destination region	Status	Actions
CrossRegionReplication	poojith-1004-test1	us-east-1	account-b-devops-bucket-050910	us-east-2	KMS-encrypted	View details, Edit rule, Delete, Actions, Create replication rule

Amazon S3 > Buckets > poojith-1004-test1

poojith-1004-test1

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Show versions

☐	Name	Type	Last modified	Size	Storage class
☐	error.html	html	September 8, 2026, 17:40:52 (UTC+05:30)	40.0 B	Standard
☐	index.html	html	September 8, 2026, 17:47:09 (UTC+05:30)	23.0 B	Standard

Amazon S3 > Buckets > account-b-devops-bucket-050910

account-b-devops-bucket-050910

Objects (4)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Show versions

☐	Name	Type	Last modified	Size	Storage class
☐	Demo file.txt	txt	September 3, 2026, 18:10:57 (UTC+05:30)	0 B	Standard
☐	error.html	html	September 8, 2026, 17:40:52 (UTC+05:30)	40.0 B	Standard
☐	index.html	html	September 8, 2026, 17:47:09 (UTC+05:30)	23.0 B	Standard
☐	job-2912662f-0b88-405e-9664-196fbb2aa191/	Folder	-	-	-

Amazon S3 > Buckets > poojith-1004-test1

poojith-1004-test1

Objects (3)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Show versions

☐	Name	Type	Last modified	Size	Storage class
☐	error.html	html	September 8, 2026, 17:40:52 (UTC+05:30)	40.0 B	Standard
☐	index.html	html	September 8, 2026, 17:47:09 (UTC+05:30)	23.0 B	Standard
☐	Screenshot 2026-09-01 175055.png	png	September 9, 2026, 12:42:26 (UTC+05:30)	211.8 KB	Standard

Amazon S3 > Buckets > account-b-devops-bucket-050910

account-b-devops-bucket-050910

Objects (5)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix



Show versions


☐	Name	Type	Last modified	Size	Storage class
☐	Demo file.txt	txt	September 3, 2026, 18:10:57 (UTC+05:30)	0 B	Standard
☐	error.html	html	September 8, 2026, 17:40:52 (UTC+05:30)	40.0 B	Standard
☐	index.html	html	September 8, 2026, 17:47:09 (UTC+05:30)	23.0 B	Standard
☐	job-2912662f-0b88-405e-9664-196fbb2aa191/	Folder	-	-	-
☐	Screenshot 2026-09-01 175055.png	png	September 9, 2026, 12:42:26 (UTC+05:30)	211.8 KB	Standard

4. Configure a bucket policy so only the Admin user can see the objects of the S3 bucket.

Implementation:

- ➔ Create an S3 bucket.
- ➔ Open the S3 bucket.
- ➔ Go to Permissions → Bucket policy.
- ➔ Add a bucket policy that allows the root user's ARN to access the bucket.
- ➔ Restrict access to the required actions such as: s3:ListBucket s3:GetObject
- ➔ Specify the root user's ARN in the Principal field.
- ➔ Save the bucket policy.
- ➔ Test using the root user.
- ➔ Verify that the root user can list and read the objects.
- ➔ Test with another user and verify that access is denied.





5. Set up lifecycle policies to automatically transition or delete objects based on specific criteria.

Implementation:

- ➔ Open the S3 bucket.
- ➔ Go to Management.
- ➔ Select Lifecycle rules.
- ➔ Click Create lifecycle rule.
- ➔ Enter a rule name.
- ➔ Choose whether the rule applies to all objects or a specific prefix/tags.
- ➔ Configure a Transition action.
- ➔ For example, move objects to S3 Standard-IA after 30 days.
- ➔ Add another transition if required, such as moving objects to S3 Glacier Flexible Retrieval after 90 days.
- ➔ Configure an Expiration action if objects should be automatically deleted.
- ➔ For example: Delete objects after 365 days
- ➔ Review the lifecycle rule.
- ➔ Click Create rule.
- ➔ Verify the rule under Management → Lifecycle rules.

aws

Search

[Alt+S]

United States (N. Virginia)

Poojith R (956035085469)

Poojith R

Amazon S3

Buckets

poojith-1004-test1

Lifecycle configuration

Create lifecycle rule

Lifecycle rule configuration

Lifecycle rule name

ArchiveOldObjects

Up to 255 characters

Choose a rule scope

☐ Limit the scope of this rule using one or more filters

☒ Apply to all objects in the bucket

⚠️

Apply to all objects in the bucket

If you want the rule to apply to specific objects, you must use a filter to identify those objects. Choose "Limit the scope of this rule using one or more filters". [Learn more](#)

☒ I acknowledge that this rule will apply to all objects in the bucket.

Lifecycle rule actions

Choose the actions you want this rule to perform.

☒ Transition current versions of objects between storage classes

This action will move current versions.

☐ Transition noncurrent versions of objects between storage classes

This action will move noncurrent versions.

☒ Expire current versions of objects

☐ Permanently delete noncurrent versions of objects

☐ Delete expired object delete markers or incomplete multipart uploads

These actions are not supported when filtering by object tags or object size.

CloudShell

Agent Toolkit for AWS

Feedback

Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

aws

Search

[Alt+S]

United States (N. Virginia)

Poojith R (956035085469)

Poojith R

Amazon S3

Buckets

poojith-1004-test1

Lifecycle configuration

Create lifecycle rule

Transition current versions of objects between storage classes

Choose transitions to move current versions of objects between storage classes based on your use case scenario and performance access requirements. These transitions start from when the objects are created and are consecutively applied. [Learn more](#)

Choose storage class transitions

Standard-IA

Days after object creation

30

Remove

Glacier Flexible Retrieval (formerly Glacier)

90

Remove

Add transition

Expire current versions of objects

For version-enabled buckets, Amazon S3 adds a delete marker and the current version of an object is retained as a noncurrent version. For non-versioned buckets, Amazon S3 permanently removes the object. [Learn more](#)

Days after object creation

365

Review transition and expiration actions

Current version actions

Noncurrent versions actions

aws

Search

[Alt+S]

United States (N. Virginia)

Poojith R (956035085469)

Poojith R

Amazon S3

Buckets

poojith-1004-test1

Lifecycle configuration

🟢 The rule "ArchiveOldObjects" has been successfully added and the lifecycle configuration has been updated

It may take some time for the configuration to be updated. Refresh the lifecycle rules list if changes to the configuration aren't displayed.

Lifecycle configuration

To manage your objects so that they are stored cost effectively throughout their lifecycle, configure their lifecycle. A lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. Lifecycle rules run once per day.

Default minimum object size for transitions

All storage classes 128K

Lifecycle rules (1)

Use lifecycle rules to define actions you want Amazon S3 to take during an object's lifetime such as transitioning objects to another storage class, archiving them, or deleting them after a specified period of time. [Learn more](#)

Find lifecycle rules by name

Lifecycle rule name	Status	Scope	Current version acti...	Noncurrent versions...	Expired object delet...	Incomplete multipa...
ArchiveOldObjects	Enabled	Entire bucket	Transition to Standard-IA, th	-	-	-

6. Push some objects to S3 using the AWS CLI.

Implementation:

- ➔ Install the AWS CLI.
- ➔ Configure AWS CLI: `aws configure`
- ➔ Enter: AWS Access Key ID AWS Secret Access Key
- ➔ Default region
- ➔ Output format
- ➔ Verify AWS CLI access: `aws sts get-caller-identity`
- ➔ List existing S3 buckets: `aws s3 ls`
- ➔ Upload a file: `aws s3 cp test.txt s3://my-devops-demo0509/`
- ➔ Verify the uploaded object: `aws s3 ls s3://my-devops-demo0509/`
- ➔ Verify the objects in the bucket.

```
poojith@poojith: ~/Downloads
$ ssh -i Test1.pem ec2-user@3.81.85.13
The authenticity of host '3.81.85.13 (3.81.85.13)' can't be established.
ED25519 key fingerprint is SHA256:Fso7HwusSkXDg0YJvt5iViLhCTZBUsoSA67a1z4V9M.
This host key is known by the following other names/addresses:
  ~/.ssh/known_hosts:52: 54.221.136.102
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '3.81.85.13' (ED25519) to the list of known hosts.

#####
##### Amazon Linux 2023
#####
##### https://aws.amazon.com/linux/amazon-linux-2023
#####

Last login: Thu Sep  3 11:25:46 2026 from 103.203.172.230
[ec2-user@ip-172-31-36-204 ~]$ sudo su -
Last login: Thu Sep  3 11:26:14 UTC 2026 on pts/1
[root@ip-172-31-36-204 ~]# exit
logout
[ec2-user@ip-172-31-36-204 ~]$ sudo su
[root@ip-172-31-36-204 ec2-user]# exit
exit
[ec2-user@ip-172-31-36-204 ~]$ sudo su -
Last login: Wed Sep  9 07:35:30 UTC 2026 on pts/1
[root@ip-172-31-36-204 ~]# aws configure
AWS Access Key ID [*****]: AKIA5G8BWSOZXPKOWS
AWS Secret Access Key [*****]: t31T0RT3pxju4qj66au5ix8wybv3E3czd+AUwb
Default region name [us-east-1]: us-east-1
Default output format [json]: json
[root@ip-172-31-36-204 ~]# aws s3 ls
2026-09-03 12:37:26 account-b-devops-bucket-050910
2026-09-08 12:08:01 poojith-1004-test1
2026-09-08 12:41:46 poojith-admin-access-test1
[root@ip-172-31-36-204 ~]# ls
test.bash
[root@ip-172-31-36-204 ~]# aws s3 cp test.bash s3://poojith-1004-test1/
upload: ./test.bash to s3://poojith-1004-test1/test.bash
[root@ip-172-31-36-204 ~]# aws s3 ls s3://poojith-1004-test1/
2026-09-09 07:12:26      216854 Screenshot 2026-09-01 175055.png
2026-09-08 12:10:52         40 error.html
2026-09-08 12:17:09          23 index.html
2026-09-09 07:39:42         362 test.bash
[root@ip-172-31-36-204 ~]#
```

7. Write a Bash script to create an S3 bucket.

- ➔ Install and configure AWS CLI.
- ➔ Create a file: `vi create-s3-bucket.sh`
- ➔ Add the following script:

```
#!/bin/bash
BUCKET_NAME="poojith-1004-test2"
REGION="us-east-1"
aws s3api create-bucket --bucket "$BUCKET_NAME" --region "$REGION"
if [ $? -eq 0 ]; then
    echo "S3 bucket $BUCKET_NAME created successfully."
else
    echo "Failed to create S3 bucket."
    exit 1

```

fi

- ➔ 4. Save the file.
- ➔ Give execute permission: `chmod +x create-s3-bucket.sh`
- ➔ Run the script: `./create-s3-bucket.sh`
- ➔ Verify the bucket: `aws s3 ls`
- ➔ Confirm that the bucket appears in the S3 bucket list.

```
poojith@ip-172-31-36-204: ~/Downloads
$ ssh -i Test1.pem ec2-user@98.80.250.218
Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023

Last login: Wed Sep  9 09:07:53 2026 from 103.203.172.230
[ec2-user@ip-172-31-36-204 ~]$ aws cli
aws: [ERROR]: argument command: Found invalid choice 'cli'

usage: aws [options] <command> [<subcommand> [<subcommand> ...] [parameters]]
To see help text, you can run:
  aws help
  aws <command> help
  aws <command> <subcommand> help

[ec2-user@ip-172-31-36-204 ~]$ aws configure
AWS Access Key ID [None]: AKIA5G8BWS0S5IGZFB
AWS Secret Access Key [None]: XwenlwyqcTcwjYYIHdn9vmt/6IAGlkp2C33pwX
Default region name [None]: us-east-1
Default output format [None]: json
[ec2-user@ip-172-31-36-204 ~]$ aws sts get-caller-identity
{
  "UserId": "956035085469",
  "Account": "956035085469",
  "Arn": "arn:aws:iam:956035085469:root"
}
```

```
[ec2-user@ip-172-31-36-204 ~]$ vi create-s3-bucket.sh
[ec2-user@ip-172-31-36-204 ~]$ chmod +x create-s3-bucket.sh
[ec2-user@ip-172-31-36-204 ~]$ ./create-s3-bucket.sh
{
  "Location": "/poojith-1004-test2",
  "BucketArn": "arn:aws:s3:::poojith-1004-test2"
}
S3 bucket poojith-1004-test2 created successfully.
[ec2-user@ip-172-31-36-204 ~]$ cat create-s3-bucket.sh
#!/bin/bash

BUCKET_NAME="poojith-1004-test2"
REGION="us-east-1"

aws s3api create-bucket --bucket "$BUCKET_NAME" --region "$REGION"

if [ $? -eq 0 ]; then
  echo "S3 bucket $BUCKET_NAME created successfully."
else
  echo "Failed to create S3 bucket."
  exit 1
fi
[ec2-user@ip-172-31-36-204 ~]$
```

The screenshot shows the AWS Management Console interface for the S3 Buckets page. The left sidebar contains navigation options like 'Amazon S3', 'Buckets', 'Files', 'Access management and security', and 'Storage management and insights'. The main content area is titled 'Buckets' and shows a list of 'General purpose buckets (4)'. The buckets listed are:

Name	AWS Region	Creation date
account-b-devops-bucket-050910	US East (Ohio) us-east-2	September 3, 2026, 18:07:26 (UTC+05:30)
poojith-1004-test1	US East (N. Virginia) us-east-1	September 8, 2026, 17:38:01 (UTC+05:30)
poojith-1004-test2	US East (N. Virginia) us-east-1	September 9, 2026, 15:36:13 (UTC+05:30)
poojith-admin-access-test1	US East (N. Virginia) us-east-1	September 8, 2026, 18:11:46 (UTC+05:30)

On the right side of the console, there are two informational panels: 'Account snapshot' (updated daily) and 'External access summary' (updated daily).

8. Upload a 1 GB file to S3 using the CLI.

Implementation:

- ➔ Make sure the AWS CLI is installed and configured.
- ➔ Verify your credentials: `aws sts get-caller-identity`
- ➔ Confirm that the destination S3 bucket exists: `aws s3 ls`
- ➔ Create or obtain a 1 GB test file.
- ➔ Upload the file using: `aws s3 cp 1gb.txt s3://poojith-1004-test2/`
- ➔ AWS CLI/S3 can use multipart upload for large objects as appropriate.
- ➔ Verify the upload: `aws s3 ls s3://poojith-1004-test2/`

```
poojith@poojith: ~/Downloads
$ aws configure
Tip: You can deliver temporary credentials to the AWS CLI using your AWS Console session by running the command 'aws login'.
AWS Access Key ID [None]: AKIA5G8BW5064F2XKQI
AWS Secret Access Key [None]: jNwB1M1c10Nuc/18FdpY8pyZ11XbdCw6eIUdujAw
Default region name [None]: us-east-1
Default output format [None]: json

poojith@poojith: ~/Downloads
$ aws sts get-caller-identity
{
  "UserId": "956035085469",
  "Account": "956035085469",
  "Arn": "arn:aws:iam:956035085469:root"
}

poojith@poojith: ~/Downloads
$ aws s3 ls
2026-09-01 18:07:26 account-b-devops-bucket-050910
2026-09-08 17:38:01 poojith-1004-test1
2026-09-08 18:11:46 poojith-admin-access-test1

poojith@poojith: ~/Downloads
$ aws s3 cp testfile_1gb.txt s3://poojith-1004-test1/
upload: ./testfile_1gb.txt to s3://poojith-1004-test1/testfile_1gb.txt

poojith@poojith: ~/Downloads
$ aws s3 ls
2026-09-01 18:07:26 account-b-devops-bucket-050910
2026-09-08 17:38:01 poojith-1004-test1
2026-09-08 18:11:46 poojith-admin-access-test1

poojith@poojith: ~/Downloads
$ aws s3 ls s3://poojith-1004-test1/
2026-09-09 12:42:26 216854 Screenshot 2026-09-01 175055.png
2026-09-08 17:40:52 40 error.html
2026-09-08 17:47:09 23 index.html
2026-09-09 13:09:42 362 test.bash
2026-09-09 15:09:50 1073741824 testfile_1gb.txt

poojith@poojith: ~/Downloads
$
```

